

INLAY: Gradient-Free Knowledge Editing with an External Addressable Memory — and Why Its Scope Mechanism Cannot Currently Be Measured

Aditya Pratap Singh

August 2026

Contents

Abstract	1
1. Introduction	2
2. Method	3
3. Experimental setup	4
4. Results: accuracy and write cost	5
5. Non-structured evidence	6
6. Multi-hop	7
7. The routing investigation, and where it ended	7
8. Why abstention cannot be measured here	9
9. Limitations	10
10. Conclusion	11
Reproducibility	11

Abstract

Deployed language models go stale the moment training ends, and retraining them costs millions. Knowledge editing asks whether a single fact can be corrected surgically instead. The dominant approach locates the weights responsible for a fact and rewrites them, which works but permanently mutates state shared with everything else the model knows.

I take the opposite position. INLAY leaves the model completely frozen and stores corrections in an external addressable memory, replaying them at decoding time by nudging the score of the intended token just before the model commits to it. Writing an edit is inserting a row: roughly five milliseconds with no gradient step, about 1600 times cheaper than gradient-based editors. Deleting an edit is exact. When no correction applies, the model’s output is unchanged bit for bit, so locality holds by construction rather than by measurement.

On CounterFact and zsRE across four model families, INLAY leads ROME, MEMIT, WISE, GRACE and AlphaEdit on edit accuracy. I extend it to AKEW, which supplies unstructured and

LLM-extracted evidence in addition to clean triples — input conditions for which weight-editing methods have no mechanism at all.

The second half of this paper reports a negative result I did not go looking for. Retrieval-based editors of this family all include a scope decision: given a query, does a stored edit apply? I built increasingly careful machinery for that decision, measured real gains from it, and then discovered the gains were not what they appeared. By executing every candidate action on 1,689 queries and recording which ones actually produced correct answers, I found that a one-line static policy is **provably optimal** on this entire benchmark suite: an oracle router that always picks the best available action ties it exactly, in all nine dataset-by-mode cells. The reason is structural. These are counterfactual benchmarks, so answering from the model’s own knowledge returns the pre-edit value and is wrong by construction. Abstention can never be correct when every question targets an edited fact.

That has a consequence beyond my own system. Scope classification is load-bearing for an entire family of architectures, and the field’s standard benchmarks cannot reward it.

1. Introduction

A model’s knowledge is not stored anywhere you can point to. There is no row in a table reading “the CEO of X is Y.” That fact exists only as a pattern distributed across billions of weights, each of which also participates in thousands of unrelated facts. This is why editing is hard, and it is the premise every method in the area has to deal with.

The dominant family — ROME, MEMIT, AlphaEdit — deals with it by locating the weights most responsible for a fact and solving for a minimal update. The localisation is real science: causal tracing corrupts the subject tokens, then restores internal activations one site at a time to find which sites recover the answer, and recovery concentrates in mid-layer MLPs at the final subject token. But every edit made this way permanently mutates weights shared with everything else, and the damage compounds under repeated editing.

I wanted to know what happens if you refuse to touch the weights at all.

1.1 Contributions

1. **INLAY**, a gradient-free editor: frozen model, external addressable memory, logit-space playback. Roughly 5 ms per write against seconds for gradient-based editors, with exact deletion and locality by construction (§2, §4).
2. **Extension to non-structured evidence.** Weight editors require a clean (**prompt**, **target**) pair and have no mechanism for raw evidence prose. I evaluate across all three AKEW input conditions (§5).
3. **A multi-hop fix and its validation.** Diagnosing why per-hop retrieval failure was


terminating chains, and a fallback that more than triples accuracy, holding at two model scales (§6).

4. **A negative result about routing, and about the benchmarks.** On 1,689 queries with per-action ground truth, an oracle router ties a one-line static policy exactly. Abstention is never uniquely correct — not once. Current benchmarks structurally cannot measure the scope decision that this architecture family depends on (§7, §8).

2. Method

Weight editing


ROME · MEMIT · AlphaEdit



each edit rewrites shared state

INLAY

frozen weights · external memory



an edit is a new row

Figure 1: Weight editing mutates state shared by every fact. INLAY leaves the model frozen and keeps corrections in a separate table consulted at decoding time. Every other property follows from this one structural choice.

2.1 The bet

Keep the model frozen. Put edits in an explicit external table. Each entry holds a key (an embedding of the fact’s question form), the answer’s token sequence, and metadata such as the relation.

Three properties fall straight out of that choice:

- **Writing is insertion.** No gradients, no optimisation, no covariance statistics to estimate. Measured at 5–15 ms.
- **Deleting is exact.** Remove the row and the edit is genuinely gone. Unlearning a gradient-installed fact is, by contrast, an open research problem — which matters when removal is a regulatory requirement rather than a nicety.
- **Edits cannot interfere.** They are physically separate state, so they cannot collide with each other or with the base model.

2.2 Addressing

Keys come from a small sentence encoder (MiniLM), chosen deliberately rather than reusing the base model’s hidden states: it is fast, its similarity geometry is purpose-trained for paraphrase matching, and it stays fixed if the base model is swapped, so one memory can in principle serve several frozen models. Keys are compressed with a Johnson–Lindenstrauss random projection.

The projection being *random* rather than learned is a deliberate trade. A learned projection would need data, would drift as the edit distribution changed, and would silently invalidate every key already stored. A fixed random matrix never invalidates the memory, at the cost of a few points of distortion against an optimal learned map. For an append-only store meant to be long-lived, that is the right side of the trade.

2.3 Playback

Retrieval always returns *something* — the nearest key. A gate decides whether that nearest match is a real hit, stacking an absolute similarity threshold, a margin check against the runner-up, and a relation-residual check that asks whether the query is about the stored *relation* rather than merely the stored subject.

Once the gate approves, INLAY must make a frozen model say the stored answer. At each decoding step the model produces a score for every vocabulary token and takes the argmax. INLAY adds a boost along the intended token’s unembedding direction so that token wins, walking the stored answer one token at a time and leaving every position outside the answer span untouched.

Intervening here rather than in hidden state is a safety argument, not a convenience. A logit boost affects exactly one token’s score. A hidden-state edit propagates through the unembedding to *every* token’s logit in ways that are hard to bound. And because the intervention is gated, when the gate does not fire the model’s behaviour is unchanged bit for bit — locality is structural, not empirical.

3. Experimental setup

Base models: GPT-J-6B, Qwen2.5-7B, Qwen2.5-1.5B-Instruct, Mistral-7B. Benchmarks: CounterFact, zsRE, RippleEdits, and AKEW (CounterFact, WikiUpdate, MQuAKE-CF across structured, unstructured and extracted input conditions).

Two methodological commitments are worth stating because they cost me results. First, every split is **subject-disjoint**: a subject never appears on both sides of a train/test boundary, which prevents a learned component from memorising an entity and appearing to generalise. Second, scoring uses one convention everywhere — a diacritic- and case-insensitive substring

match against gold and its aliases — rather than each method’s internal metric, so numbers are comparable across methods that report differently.

Baselines run through EasyEdit. One implementation note that materially affects correctness: `BaseEditor.edit(sequential_edit=False)` restores weights before returning, so generating after it returns silently scores the *unedited* model. Every baseline here uses `sequential_edit=True` with an explicit state-dict snapshot and restore under harness control.

4. Results: accuracy and write cost




Figure 2: Edit accuracy on CounterFact, structured input, GPT-J-6B, n=147, scored with the shared convention rather than each method’s internal metric.

INLAY reaches 100% on this cell. AlphaEdit is the strongest weight editor at 89.12%, consistent with its design: constraining the update to the null space of preserved knowledge makes each edit more surgical. ROME and MEMIT land at an identical 83.67% — plausible rather than suspicious, since single-fact CounterFact is where both converge and MEMIT’s real advantage is multi-edit preservation, not single-edit accuracy.

Two results in that figure deserve their own sentence, because both are failures of methods that work fine elsewhere.

WISE reaches only 66.67% despite its edits landing. EasyEdit’s own post-edit metric reports success on 141 of 147 edits, yet generation-time accuracy is far below ROME’s. WISE routes through a side-memory module at generation time, and that routing is an additional point of failure that in-place weight updates do not have. The edit registers; it is not reliably retrieved.

GRACE does not land at all. Its post-edit accuracy reads exactly 0.0 on every edit, correctly triggering a fail-loud guard rather than reporting a misleading number. This agrees with an independent generation-based harness I ran earlier: GRACE’s radius-based codebook matching

almost never fires on paraphrases. Two independent harnesses agreeing is why I record this as a real limitation on this model and dataset rather than a bug I failed to chase down.




Figure 3: Write cost per edit. INLAY’s figure is measured directly; the weight-editing figure is the aggregate ratio this project reports, plotted as one band rather than fabricated per method.

The write-cost gap is the systems argument, and it is not marginal. Inserting a row is not a cheaper optimisation — it is the absence of one.

5. Non-structured evidence

Weight-editing methods need a (**prompt**, **target**) pair to compute their update from. They have no mechanism for “here is a paragraph of evidence, extract and install whatever matters.” Building that extraction step is a separate research problem, so on AKEW’s unstructured and extracted conditions the weight-editing family is not merely worse — it is inapplicable. I state this as a structural limitation rather than working around it.

INLAY’s retrieval-based design handles all three conditions. Accuracy is lower on extracted input (78.23% on CounterFact) than on unstructured prose (87.07%), which traces to extraction noise rather than retrieval failure: retrieval finds the right card 98.64% of the time either way. An LLM-extracted triple is a lossier summary than the prose it came from.

While running this I found a real bug in my own router. It was selecting the recite-directly path for all 147 unstructured queries, even though that path degrades to reciting a raw evidence sentence outside structured mode. Gating the decision on input mode recovered the gap: unstructured moved from 85.71% to 87.07%, and structured stayed at 100%. It was invisible to every oracle-evidence experiment I had run, because those never exercised the live routing decision.




Figure 4: Iterative decomposition versus naive single-shot on MQuAKE-CF, after the per-hop fallback fix.

6. Multi-hop

My first iterative multi-hop loop scored 5.0% against a naive single-shot baseline’s 22.5% — more than four times *worse* than doing nothing clever. Every sampled failure produced no answer at all, which pointed at termination rather than accuracy.

The cause was a wrong assumption. MQuAKE-CF’s chains mix edited facts with ordinary unedited world facts: 277 of 354 groups contain exactly one edit across a two- or three-hop chain. My loop treated every hop as something that must be retrieved and verified against the edit index, so when a hop needed a fact that was never edited — and therefore was not in the index — retrieval correctly found nothing, and the chain terminated answerless.

The fix is to treat a retrieval miss as a signal to fall back to the model’s own knowledge for that hop and **continue**, rather than as chain failure. That took the loop from 5.0% to 47.5% on the same sample, and it holds at scale: 53.95% against 16.10% on the full 354-group pool, and 42.00% against 20.67% at 7B.

One failure pattern recurred identically across all three runs: two genuinely different questions, both routing through a “founder of a religion” hop, converging on the same wrong answer. Three independent runs reproducing the same specific error is evidence of a real confusion rather than noise.

7. The routing investigation, and where it ended

Everything above rests on a scope decision: given a query, does a stored edit apply, and what should be done about it? I spent the largest single block of this project on that decision, and the result is not what I expected.

7.1 The failure that started it

On MQuAKE-CF, gating was *net-negative in every input mode*: the routed pipeline underperformed simply disabling the gates and always reasoning, by up to 15.87 points. Two natural fixes failed, and their failure was diagnostic. Recalibrating the verifier threshold moved the false-fire rate from 18.33% to 18.22% — no meaningful change, because the score distributions genuinely overlap. Raising the confidence threshold from 0.85 to 0.97 changed *nothing*: byte-identical decisions, because confidence on the *wrong* retrievals was already above 0.97.

The router was asking how confident the verifier was, when what it needed to know was whether that confidence was *discriminative*. A verifier scoring 0.99 on the top candidate and 0.98 on four unrelated ones is not confident; it is saturated. No threshold on a top-1 score can see that.

7.2 A second-order signal, which worked

So I built a head that predicts retrieval reliability from the *shape* of the candidate set rather than the magnitude of its top score: margins to the next best candidate, how many candidates clear the threshold, entropy over the neighbourhood, subject diversity among the top-k. Trained on CounterFact and WikiUpdate with MQuAKE-CF held out entirely, it reached **0.956 AUROC on the held-out dataset it had never seen**.

The mechanism showed up in the fitted weights rather than being asserted: 71.5% of coefficient mass sat on margin features, and the single largest weight was -1.43 on the best *competing* candidate’s score. The more confident the verifier is about a rival card, the less trustworthy its top pick. That is exactly the quantity a first-order threshold cannot express.

Used to gate the router adaptively, it produced real, statistically significant gains: **+15.9 points on two MQuAKE-CF cells ($p = 0.0019$, Holm-corrected)**, while making byte-identical decisions — zero discordant pairs — on the cells where fixed gating already worked. The gain was identical at 1.5B and 7B.

I thought this was the contribution.

7.3 The result that reframed it

Two follow-up experiments failed in ways that pointed the same direction. A three-way policy using the *magnitude* of predicted unreliability failed on the cell it was designed to fix and was catastrophic on MQuAKE-CF. Adding the head to the multi-hop per-hop gate changed nothing at all.

So I stopped improving the router and measured the ceiling instead. For 1,689 queries spanning all three datasets and all three input conditions, I executed *every* candidate action and scored the result, giving per-query ground truth for which actions actually work.

An oracle router that picks the best available action on every query is exactly equal to a one-line static policy — “recite directly where that is legal, otherwise reason over the




Figure 5: Every candidate action executed and scored on 1,689 queries. The oracle marker sits exactly on the static-policy bar in all nine cells.

retrieved evidence.” Identical to four decimal places in all nine cells. The maximum possible gain of any per-query routing method, learned or otherwise, over one line of code is **0.00 points**.

The mechanism is stark. Abstention succeeded on 19 of 1,689 queries, and on all 19 of those, reasoning or direct recitation succeeded too. **Abstention is the sole winning action zero times.**

7.4 What that does to §7.2

The measured gains in §7.2 are real and I stand behind the numbers. The *explanation* was wrong. I framed the head as adapting its decision per query to the local regime. It does not, because there is no per-query decision worth making. What it actually does is suppress gates that misfire — that is, it converges toward the static policy. Its gains measure the harm the fixed gates were doing, not insight it contributes.

A one-line rule matches it everywhere, with no training data, no out-of-distribution generalisation argument, and none of the fivefold cross-encoder cost. I would rather report that than let a well-measured number carry a wrong story.

8. Why abstention cannot be measured here

The reason is structural and, once seen, obvious.

These are **counterfactual** benchmarks. The evaluation question asks for the *post-edit* answer. Answering from parametric knowledge alone returns the *pre-edit* value — which is, by construction, wrong. There is no query in the suite for which “no edit applies, answer from what you already know” is the correct behaviour.

Abstention cannot be correct on a benchmark where every question targets an edited fact.

This generalises past my system. Scope classification is load-bearing for the SERAC-descendant design family: the entire premise is a learned decision about whether a stored edit applies. But on these benchmarks that decision is degenerate — the answer is always “yes.” **A benchmark with no negatives cannot measure a classifier’s ability to reject**, and any abstention path evaluated on one can only ever lose points.

It also explains something about my own system that had been sitting unexplained. On RippleEdits, which *does* contain same-subject different-relation queries and therefore has partial negatives, INLAY’s preservation score was the worst in the comparison at 0.05 — keys over-firing on queries about the right subject but the wrong relation. Scope precision matters exactly where a benchmark has negatives, and that is precisely where INLAY is weakest and where the main benchmarks are silent.

The missing evaluation condition is constructible: ask evaluation questions against an index that does *not* contain the corresponding edit. On such queries abstention becomes correct, the oracle should pull away from the static policy, and routing acquires real headroom for the first time. That is the experiment I would run next, and it measures what every scope-classifier paper claims to measure.

9. Limitations

Recitation, not reasoning. Logit-space playback reproduces a stored answer. It cannot combine that answer with anything else, because by the final decoding step there is no reasoning left to influence. This is structural, and it is why compositional queries need the retrieve-into-context path rather than playback.

Scope precision on hard negatives. The 0.05 RippleEdits preservation score is the real open problem, and §8 explains why the field’s main benchmarks did not force me to confront it earlier.

Small cells. The MQuAKE-CF slices are $n=63$. The +6.4-point structured-mode result rests on four discordant pairs, where the smallest attainable exact McNemar p-value is 0.125 — suggestive, not established, and reported that way.

One retrieval stack. A single encoder and one cross-encoder verifier throughout; sensitivity to those choices is untested.

A scale anomaly I diagnosed but did not fix. WikiUpdate is 4.38 points *worse* at 7B than at 1.5B. It is not an editing failure: the larger model refuses to answer far more often on noisy evidence (46.88% versus 29.38% overall, 91.11% versus 66.67% where retrieval is wrong), and substring accuracy scores a refusal identically to a wrong answer. The larger model is more accurate whenever it *commits* — it simply commits less. This is a property of the metric, not of any method, and it means scale comparisons on noisy-retrieval benchmarks will systematically understate larger models.

10. Conclusion

Keeping the weights frozen and putting edits in an external addressable memory buys three things that gradient-based editing cannot offer: writes about 1600 times cheaper, exact deletion, and locality that holds by construction. It pays for them by reciting facts rather than understanding them, which is a real cost and the honest boundary of the approach.

The part I did not expect to write is §7 and §8. I built a scope mechanism, measured genuine gains from it, and then found that the benchmarks could not distinguish those gains from repairing damage that a simpler system would never have caused. A one-line policy is provably optimal on all 1,689 queries I tested, because abstention is never the right answer when every question targets an edited fact.

I think that is the more useful finding. It suggests that reported routing gains on these benchmarks measure how much harm a gate was doing, and that an entire architectural family is being evaluated on suites that cannot reward its central component. The fix is not a better router. It is a benchmark that contains queries where the right answer is “this edit does not apply.”

Reproducibility

All numbers come from logged runs in this project’s repository. Splits are subject-disjoint with seed 0; the held-out dataset is never touched during training of any learned component.

<code>src/akew_outcome_labels.py</code>	per-action ground truth (§7.3)
<code>src/akew_headroom.py</code>	the oracle-versus-static table
<code>src/akew_reliability.py</code>	the second-order head (§7.2)
<code>src/akew_stats.py</code>	Wilson intervals, exact McNemar, paired bootstrap, Holm correction
<code>src/akew_eval_weightedit.py</code>	ROME / MEMIT / AlphaEdit / WISE / GRACE
<code>paper/make_figures.py</code>	every figure in this paper

Self-tests accompany the statistical and learned components (55 assertions across the reliability,

outcome-head and statistics modules), including checks that the paired bootstrap preserves pairing and that a reordered feature checkpoint refuses to load rather than silently mispredicting.